

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: from sklearn.preprocessing import Imputer  
from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LinearRegression  
from sklearn.tree import DecisionTreeRegressor  
from sklearn.ensemble import RandomForestRegressor  
  
from sklearn import metrics
```

```
In [3]: housing = pd.read_csv('housing1.csv')  
housing
```

Out[3]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	household
0	-122.23	37.88	41	880	129.0	322	
1	-122.22	37.86	21	7099	1106.0	2401	
2	-122.24	37.85	52	1467	190.0	496	
3	-122.25	37.85	52	1274	235.0	558	
4	-122.25	37.85	52	1627	280.0	565	
5	-122.25	37.85	52	919	213.0	413	
6	-122.25	37.84	52	2535	489.0	1094	
7	-122.25	37.84	52	3104	687.0	1157	
8	-122.26	37.84	42	2555	665.0	1206	
9	-122.25	37.84	52	3549	707.0	1551	
10	-122.26	37.85	52	2202	434.0	910	
11	-122.26	37.85	52	3503	752.0	1504	
12	-122.26	37.85	52	2491	474.0	1098	
13	-122.26	37.84	52	696	191.0	345	
14	-122.26	37.85	52	2643	626.0	1212	
15	-122.26	37.85	50	1120	283.0	697	
16	-122.27	37.85	52	1966	347.0	793	
17	-122.27	37.85	52	1228	293.0	648	
18	-122.26	37.84	50	2239	455.0	990	
19	-122.27	37.84	52	1503	298.0	690	
20	-122.27	37.85	40	751	184.0	409	
21	-122.27	37.85	42	1639	367.0	929	
22	-122.27	37.84	52	2436	541.0	1015	
23	-122.27	37.84	52	1688	337.0	853	
24	-122.27	37.84	52	2224	437.0	1006	
25	-122.28	37.85	41	535	123.0	317	
26	-122.28	37.85	49	1130	244.0	607	
27	-122.28	37.85	52	1898	421.0	1102	
28	-122.28	37.84	50	2082	492.0	1131	
29	-122.28	37.84	52	729	160.0	395	
...	...	...	...	...	...	...	...
20610	-121.56	39.10	28	2130	484.0	1195	
20611	-121.55	39.10	27	1783	441.0	1163	
20612	-121.56	39.08	26	1377	289.0	761	
20613	-121.55	39.09	31	1728	365.0	1167	

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	household
<b>20614</b>	-121.54	39.08	26	2276	460.0	1455	
<b>20615</b>	-121.54	39.08	23	1076	216.0	724	
<b>20616</b>	-121.53	39.08	15	1810	441.0	1157	
<b>20617</b>	-121.53	39.06	20	561	109.0	308	
<b>20618</b>	-121.55	39.06	25	1332	247.0	726	
<b>20619</b>	-121.56	39.01	22	1891	340.0	1023	
<b>20620</b>	-121.48	39.05	40	198	41.0	151	
<b>20621</b>	-121.47	39.01	37	1244	247.0	484	
<b>20622</b>	-121.44	39.00	20	755	147.0	457	
<b>20623</b>	-121.37	39.03	32	1158	244.0	598	
<b>20624</b>	-121.41	39.04	16	1698	300.0	731	
<b>20625</b>	-121.52	39.12	37	102	17.0	29	
<b>20626</b>	-121.43	39.18	36	1124	184.0	504	
<b>20627</b>	-121.32	39.13	5	358	65.0	169	
<b>20628</b>	-121.48	39.10	19	2043	421.0	1018	
<b>20629</b>	-121.39	39.12	28	10035	1856.0	6912	
<b>20630</b>	-121.32	39.29	11	2640	505.0	1257	
<b>20631</b>	-121.40	39.33	15	2655	493.0	1200	
<b>20632</b>	-121.45	39.26	15	2319	416.0	1047	
<b>20633</b>	-121.53	39.19	27	2080	412.0	1082	
<b>20634</b>	-121.56	39.27	28	2332	395.0	1041	
<b>20635</b>	-121.09	39.48	25	1665	374.0	845	
<b>20636</b>	-121.21	39.49	18	697	150.0	356	
<b>20637</b>	-121.22	39.43	17	2254	485.0	1007	
<b>20638</b>	-121.32	39.43	18	1860	409.0	741	
<b>20639</b>	-121.24	39.37	16	2785	616.0	1387	

20640 rows × 10 columns



In [4]: `housing.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude                20640 non-null float64
latitude                 20640 non-null float64
housing_median_age       20640 non-null int64
total_rooms              20640 non-null int64
total_bedrooms           20433 non-null float64
population               20640 non-null int64
households               20640 non-null int64
median_income            20640 non-null float64
ocean_proximity          20640 non-null object
median_house_value       20640 non-null int64
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB
```

In [5]: `print(housing.head())`

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	

  

	population	households	median_income	ocean_proximity	median_house_value
0	322	126	8.3252	NEAR BAY	452600
1	2401	1138	8.3014	NEAR BAY	358500
2	496	177	7.2574	NEAR BAY	352100
3	558	219	5.6431	NEAR BAY	341300
4	565	259	3.8462	NEAR BAY	342200

In [6]: `housing.isnull().sum()`

```
Out[6]: longitude                0
latitude                       0
housing_median_age             0
total_rooms                    0
total_bedrooms                 207
population                     0
households                     0
median_income                  0
ocean_proximity                0
median_house_value             0
dtype: int64
```

```
In [7]: print ("Total_bedrooms column Mode is "+str(housing["total_bedrooms"].mode())
        +"\n")
        print(housing["total_bedrooms"].describe())
```

Total\_bedrooms column Mode is 0 280.0  
dtype: float64

```
count    20433.000000
mean      537.870553
std       421.385070
min        1.000000
25%       296.000000
50%       435.000000
75%       647.000000
max      6445.000000
Name: total_bedrooms, dtype: float64
```

```
In [8]: print(housing.iloc[:,4:5].head())
        imputer = Imputer(np.nan, strategy="median")
        imputer.fit(housing.iloc[:,4:5])
        housing.iloc[:,4:5] = imputer.transform(housing.iloc[:,4:5])
        housing.isnull().sum()
```

```
total_bedrooms
0          129.0
1         1106.0
2          190.0
3          235.0
4          280.0
```

E:\New folder (2)\lib\site-packages\sklearn\utils\deprecation.py:58: DeprecationWarning: Class Imputer is deprecated; Imputer was deprecated in version 0.20 and will be removed in 0.22. Import impute.SimpleImputer from sklearn instead.

```
warnings.warn(msg, category=DeprecationWarning)
```

```
Out[8]: longitude      0
        latitude       0
        housing_median_age  0
        total_rooms     0
        total_bedrooms  0
        population     0
        households     0
        median_income   0
        ocean_proximity  0
        median_house_value  0
        dtype: int64
```

```
In [9]: labelEncoder = LabelEncoder()
print(housing["ocean_proximity"].value_counts())
housing["ocean_proximity"] = labelEncoder.fit_transform(housing["ocean_proximity"])
housing["ocean_proximity"].value_counts()
housing.describe()
```

```
<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
```

```
Name: ocean_proximity, dtype: int64
```

Out[9]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	popul
<b>count</b>	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.00
<b>mean</b>	-119.569704	35.631861	28.639486	2635.763081	536.838857	1425.47
<b>std</b>	2.003532	2.135952	12.585558	2181.615252	419.391878	1132.46
<b>min</b>	-124.350000	32.540000	1.000000	2.000000	1.000000	3.00
<b>25%</b>	-121.800000	33.930000	18.000000	1447.750000	297.000000	787.00
<b>50%</b>	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.00
<b>75%</b>	-118.010000	37.710000	37.000000	3148.000000	643.250000	1725.00
<b>max</b>	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.00

```
In [11]: housing_ind = housing.drop("median_house_value",axis=1)
print(housing_ind.head())
housing_dep = housing["median_house_value"]
print("Medain Housing Values")
print(housing_dep.head())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	

	population	households	median_income	ocean_proximity
0	322	126	8.3252	3
1	2401	1138	8.3014	3
2	496	177	7.2574	3
3	558	219	5.6431	3
4	565	259	3.8462	3

Medain Housing Values

```
0    452600
1    358500
2    352100
3    341300
4    342200
```

Name: median\_house\_value, dtype: int64

```
In [12]: #check for rand_state
X_train,X_test,y_train,y_test = train_test_split(housing_ind,housing_dep,test_
size=0.2,random_state=42)
#print(X_train.head())
#print(X_test.head())
#print(y_train.head())
#print(y_test.head())
print("X_train shape {} and size {}".format(X_train.shape,X_train.size))
print("X_test shape {} and size {}".format(X_test.shape,X_test.size))
print("y_train shape {} and size {}".format(y_train.shape,y_train.size))
print("y_test shape {} and size {}".format(y_test.shape,y_test.size))
```

X\_train shape (16512, 9) and size 148608

X\_test shape (4128, 9) and size 37152

y\_train shape (16512,) and size 16512

y\_test shape (4128,) and size 4128



```
In [13]: X_train.head()
```

Out[13]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	household
<b>14196</b>	-117.03	32.71	33	3126	627.0	2300	
<b>8267</b>	-118.16	33.77	49	3382	787.0	1314	
<b>17445</b>	-120.48	34.66	4	1897	331.0	915	
<b>14265</b>	-117.11	32.69	36	1421	367.0	1418	
<b>2271</b>	-119.80	36.78	43	2382	431.0	874	

```
In [14]: independent_scaler = StandardScaler()
X_train = independent_scaler.fit_transform(X_train)
X_test = independent_scaler.transform(X_test)
print(X_train[0:5,:])
print("test data")
print(X_test[0:5,:])
```

E:\New folder (2)\lib\site-packages\sklearn\preprocessing\data.py:625: DataConversionWarning: Data with input dtype int32, int64, float64 were all converted to float64 by StandardScaler.

```
return self.partial_fit(X, y)
```

```
[[ 1.27258656 -1.3728112  0.34849025  0.22256942  0.21122752  0.76827628
  0.32290591 -0.326196  2.00593172]
 [ 0.70916212 -0.87669601  1.61811813  0.34029326  0.59309419 -0.09890135
  0.6720272  -0.03584338  2.00593172]
 [-0.44760309 -0.46014647 -1.95271028 -0.34259695 -0.49522582 -0.44981806
 -0.43046109  0.14470145  2.00593172]
 [ 1.23269811 -1.38217186  0.58654547 -0.56148971 -0.40930582 -0.00743434
 -0.38058662 -1.01786438  2.00593172]
 [-0.10855122  0.5320839  1.14200767 -0.11956547 -0.25655915 -0.48587717
 -0.31496232 -0.17148831 -0.1124266  ]]
```

test data

```
[[ 0.28534728  0.1951  -0.28632369 -0.52286157 -0.24701249 -0.03030109
 -0.37008673 -1.15508475 -0.1124266  ]
 [ 0.06097472 -0.23549054  0.11043502  0.13841528 -0.24701249  0.12185077
  0.220532  -0.70865905 -0.1124266  ]
 [-1.42487026  1.00947776  1.85617335  0.54630997 -0.24701249 -0.10241931
  1.21539643 -0.21040155  1.29981228]
 [ 0.42994293 -0.63799909 -0.92113763  0.18808002 -0.24701249  0.24497944
 -0.01309052  0.97511311 -0.81854604]
 [-1.17058135  0.45719859  0.42784199 -0.13382109 -0.24701249 -0.31965346
 -0.18896365 -0.08179356  2.00593172]]
```

E:\New folder (2)\lib\site-packages\sklearn\base.py:462: DataConversionWarning: Data with input dtype int32, int64, float64 were all converted to float64 by StandardScaler.

```
return self.fit(X, **fit_params).transform(X)
```

E:\New folder (2)\lib\site-packages\ipykernel\_launcher.py:3: DataConversionWarning: Data with input dtype int32, int64, float64 were all converted to float64 by StandardScaler.

This is separate from the ipykernel package so we can avoid doing imports until

```
In [15]: #initantiate the linear regression
linearRegModel = LinearRegression(n_jobs=-1)
#fit the model to the training data (learn the coefficients)
linearRegModel.fit(X_train,y_train)
#print the intercept and coefficients
print("Intercept is "+str(linearRegModel.intercept_))
print("coefficients is "+str(linearRegModel.coef_))
```

```
Intercept is 207194.69373788778
coefficients is [-85854.94724101 -90946.06271148  14924.30655143 -17693.2340
5277
 48767.60670995 -43884.16852449  17601.31495096  77144.10164179
-451.52015229]
```

```
In [16]: y_pred = linearRegModel.predict(X_test)
```

```
In [17]: print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
print(np.sqrt(metrics.mean_squared_error(y_train,linearRegModel.predict(X_train))))
```

```
71147.87146118373
69361.0714290645
```

```
In [18]: dtReg = DecisionTreeRegressor(max_depth=9)
dtReg.fit(X_train,y_train)
```

```
Out[18]: DecisionTreeRegressor(criterion='mse', max_depth=9, max_features=None,
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
presort=False, random_state=None, splitter='best')
```

```
In [19]: dtReg_y_pred = dtReg.predict(X_test)
dtReg_y_pred
```

```
Out[19]: array([ 60503.2556391 ,  75919.52054795, 478283.56097561, ...,
488611.25          ,  75919.52054795, 211563.96963563])
```

```
In [20]: print(np.sqrt(metrics.mean_squared_error(y_test,dtReg_y_pred)))

60981.93860192401
```

```
In [21]: rfReg = RandomForestRegressor(30)
rfReg.fit(X_train,y_train)
```

```
Out[21]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=30, n_jobs=None,
oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
In [23]: rfReg_y_pred = rfReg.predict(X_test)
print(len(rfReg_y_pred))
print(len(y_test))
print(rfReg_y_pred[0:5])
print(y_test[0:5])
```

```
4128
4128
[ 48646.66666667  71933.33333333  466903.7          278603.33333333
 245740.         ]
20046      47700
3024      45800
15663     500001
20484     218600
9814      278000
Name: median_house_value, dtype: int64
```

```
In [24]: print(np.sqrt(metrics.mean_squared_error(y_test, rfReg_y_pred)))

50631.69730545329
```

```
In [ ]:
```